

Positioning: ARES-Harness vs. SOTA injection defenses

System	Surface guarded	Mechanism	Conclusion-integrity?
Dual-LLM (Willison '23)	input (ctrl/data sep)	quarantined LLM	X
Spotlighting (Hines '24)	input	prompt marking (model)	X
Instr. Hierarchy (Wallace '24)	input priority	learned (training)	X
StruQ (Chen '24)	input	struct queries + learned	X
SecAlign (Chen '24)	input	preference-opt (learned)	X
CaMeL (Debenedetti '25)	action	cap interp-IFC (det.)	X
ARES-Harness (this work)	input + action	det. provenance-taint (no LLM in decision)	✓

← *the decision-integrity axis — orthogonal to all surface defenses; only ARES scores it*